

world money__ball - kaggle dataset

patrick mcdevitt (& collaboration with preeti swaminathan)

04 juin 2017

money__ball

example of ordinary logistic regression model building techniques

1. DATA EXPLORATION (40 points)

This data set contains approximately 2200 records. Each record represents a professional baseball team from the years 1871 to 2006 inclusive. Each record has the performance of the team for the given year, with all of the statistics adjusted to match the performance of a 162 game season.

Name	_ext	Number	Mean	St Dev	Min	Max
index		2,246	1,265,817	734,970	1	2,534
target_wins		2,246	81,178	14,596	39	123
team_batting	_h	2,238	1,465,629	128,777	992	2,192
team_batting	_2b	2,246	241,694	46,136	113	458
team_batting	_3b	2,243	54,875	27,273	0	190
team_batting	_hr	2,246	100,715	60,156	0	264
team_batting	_bb	2,246	505,620	116,931	45	878
team_batting	_so	2,145	741,267	242,532	0	1,399
team_baserun	_sb	2,126	123,686	84,771	18	562
team_baserun	_cs	1,490	52,299	20,955	11	166
team_batting	_hbp	191	59,356	12,967	29	95
team_pitching	_h	2,234	1,650,051	548,521	1,137	7,093
team_pitching	_hr	2,246	106,649	61,037	0	343
team_pitching	_bb	2,238	547,450	114,549	119	1,123
team_pitching	_so	2,141	801,531	252,156	0	3,450
team_fielding	_e	2,240	232,420	192,012	65	1,374
team_fielding	_dp	1,983	146,538	26,123	52	228
team_batting	_1b	2,240	1,069,086	116,599	709	1,706
team_batting	_ab	2,238	5,788,027	130,347	5,317	6,517
team_batting	_slg	2,246	0,366	0,038	0,244	0,481

Name	_ext	Number	Mean	St Dev	Min	Max
team_batting	_obp	2,245	0,311	0,019	0,235	0,392

Describe the size and the variables in the MONEYBALL data set so that a manager can understand it. Consider that too much detail will cause a manager to lose interest while too little detail will make the manager consider that you aren't doing your job. Some suggestions are given below. Please do NOT treat this as a check list of things to do to complete the assignment. You should have your own thoughts on what to tell the boss. These are just ideas.

- Mean / Standard Deviation / Median
- Bar Chart or Box Plot of the data
- Is the data correlated to the target variable (or to other variables?)
- Are any of the variables missing and need to be imputed "fixed"?

2. DATA PREPARATION (40 Points)

Each variable was assessed for several relevant characteristics to decide potential improvements for incorporation (or not) into the regression model. * what is the percentage of missing values in that column ? * is there a strong correlation between each column and the other potential independent variables ? * what is the single parameter correlation to the dependent variable (target wins) ? * what is the skewness for each variable ? * is there a generally linear relationship to the target_wins ? * based on a linear single parameter regression fit, are the residuals from this fit uniformly distributed ?

All of the above questions resulted in making the following decisions / actions on the data set :

- team_batting_hbp - dropped due to > 90% missing values ... little likelihood to make an informed and useful imputation correction on that data column
- team_pitching_hr - dropped due to 97% correlation with team_batting_hr. There is no obvious explanation for this. Perhaps there is a mistake in the data tabulation that mixed pitching home runs allowed with batting home runs achieved or perhaps there is a tim. There is no obvious reason to expect that a team's offense that hits a lot (or a few) home runs that their defense (pitching staff) also allows (or restricts) home runs to approximately the same level. In any case, for two independent characteristics that are 97% correlated, it is not value added to include both characteristics in the regression model.
- team_batting_h (hits) - decomposed to create new predictor "team_batting_1b" (singles), with the idea that doubles (_2b), triples (_3b), and home runs (_hr) are all a part of "hits", it is perhaps useful to separate out the potential influence of singles (_1b) from the other hits, and drop hits (_h) from the independent predictors data set. There is, at the least, no loss of information from this choice.
- team_batting_slg - new predictor added (created) from some of the existing columns. "Slugging Percentage" is traditionally a high value predictor in SABR metrics evaluations of win predictions : http://www.baseball-reference.com/bullpen/Slugging_percentage. Although all of the characteristics that constitute the generally accepted definition of slugging percentage, (at-bats is not a part of the data set), an estimate of at-bats can be inferred based on hits, outs produced from on-base runners, and minimum number of at-bats in a game.
- team_batting_obp - new predictor added. Similar to "slugging percentage", "on base percentage" is another widely metric in SABR analyses for win predictions, http://www.baseball-reference.com/bullpen/On_base_percentage.

Describe how you have transformed the data by changing the original variables or creating new variables. If you did transform the data or create new variables, discuss why you did this. Here are some possible transformations.

- a. Fix missing values (maybe with a Mean or Median value or use a decision tree)
- b. Create flags to suggest if a variable was missing.
- c. Transform data by putting it into buckets
- d. Mathematical transforms such as log or square root
- e. Combine variables (such as ratios or adding or multiplying) to create new variables

3. BUILD MODELS (40 Points)

Build at least three different LINEAR REGRESSION using different variables (or the same variables with different transformations). You may select the variables manually, use an approach such as Forward or Stepwise, use a different approach such as trees, or use a combination of techniques. Describe the techniques you used. If you manually selected a variable for inclusion into the model or exclusion into the model, indicate why this was done.

Discuss the coefficients in the model, do they make sense? For example, if a team hits a lot of Home Runs, it would be reasonably expected that such a team would win more games. However, if the coefficient is negative (suggesting that the team would lose more games), then that needs to be discussed. Are you keeping the model even though it is counter intuitive? Why? The boss needs to know.

4. SELECT MODELS (40 Points)

Decide on the criteria for selecting the “Best Model”. Will you use a metric such as Adjusted R-Square or AIC? Will you select a model with slightly worse performance if it makes more sense or is more parsimonious? Discuss why you selected your model. STAND ALONE SCORING PROGRAM (40 POINTS)

WRITE MODEL DEPLOYMENT CODE (40 Points)

Write a Stand Alone SAS data step that will score new data and predict the number of wins. The variable with the Predicted number of Wins should be named:

P_TARGET_WINS

The SAS data step will need to include:

- a. All the variable transformations such as fixing missing values
- b. The regression formula

SCORED DATA FILE (50 POINTS)

SCORE THE MONEYBALL_TEST DATA SET (50 Points)

Use the stand alone program that you wrote in the previous section. Score the data file MONEYBALL_TEST. Create a file that has only TWO variables for each record:

INDEX P_TARGET_WINS

The first variable, INDEX, will allow me to match my grading key to your predicted value. If I cannot do this, you won't get a grade. So please include this value. The second value, P_TARGET_WINS is the number of wins you believe the team will have in season based upon the data given to you.

Your values will be compared against ... A Perfect Model Instructor's Model Performance of Other Students Predict the Average value for everybody (MEAN) Random Model Worst Possible Model

If your model is not better than simply using an AVERAGE value, you will receive negative points If your model is not better than generating a RANDOM value, you will receive a LOT of negative points If your model is not better than the WORST model, then it will be a WHOLE LOT of negative points.